

What Can IRT Person-Fit Statistics Reveal About SFT-Stage Benchmark Contamination?

Anonymous ACL submission

Abstract

Benchmark contamination in Large Language Models (LLMs) undermines evaluation validity, yet most detection methods require privileged access to training data, model internals, or output probabilities. We present the first controlled characterization of Item Response Theory (IRT) person-fit statistics for detecting supervised fine-tuning (SFT) contamination, requiring only binary response patterns and calibrated four-parameter logistic item parameters; diagnostic θ -anchoring requires a clean baseline. A contaminated LLM produces aberrant response patterns analogous to a test-taker with item preknowledge. Using ℓ_z^* in a matched-difference design ($\Delta\ell_z^*$) across three 7–8B model families ($n=10$ clean seeds per family) on MMLU (14,042 items) and ARC-Challenge (295 items), we identify three boundaries: (1) model-specific θ -shift dynamics, where contamination-induced ability inflation absorbs the person-fit signal; θ -anchoring amplifies detection by $2.4\times$ while reducing coefficient of variation from 60% to 19%; (2) a detection-diagnosis tradeoff where simpler statistics yield stronger detection but lack ℓ_z^* 's diagnostic capabilities (seen/unseen decomposition, mechanistic θ -shift analysis); and (3, exploratory) ability-dependent non-monotonic dose-response under free- θ estimation, driven by θ -shift absorption; θ -anchored analysis reveals stable positive dose-response. These findings establish when and why binary response patterns carry contamination information beyond aggregate accuracy¹.

1 Introduction

Benchmark contamination (the unintended inclusion of evaluation data in training corpora) poses a growing threat to LLM evaluation integrity (Dodge et al., 2021; Magar and Schwartz, 2022; Sainz et al.,

2023), as the probability of overlap increases structurally with expanding corpora and benchmarks (Jacovi et al., 2023; Chang et al., 2024). Existing detection methods all require privileged access beyond binary response records: training corpora, token-level probabilities, crafted prompts, or multiple test variants (§2). Moreover, all log-probability-based detectors fail after RL post-training (Wang et al., 2025), motivating the question: *what can binary response patterns alone reveal about contamination?*

A contaminated LLM is functionally analogous to a test-taker with item preknowledge: both produce response patterns where difficult items are answered correctly through memorization rather than genuine ability (Figure 1a–b). This connection points to person-fit statistics from psychometric cheating detection (Drasgow et al., 1985; Snijders, 2001; Karabatsos, 2003). We focus on SFT-stage contamination, where benchmark answers are explicitly included in fine-tuning data, as it provides controlled ground truth and is increasingly relevant as fine-tuning-as-a-service platforms proliferate (Figure 2).

Using ℓ_z^* in a matched-difference design across three 7–8B model families (Qwen2.5-7B, Llama-3.1-8B, Mistral-7B; $n=10$ clean seeds each) with pre-calibrated four-parameter logistic (4PL) item parameters (Zhou et al., 2026) on MMLU (Hendrycks et al., 2021) and ARC-Challenge, we identify three behavioral boundaries:

1. **Model-specific θ -shift dynamics.** Contamination inflates ability estimates ($\hat{\theta}$), absorbing the person-fit signal; θ -anchoring amplifies Qwen’s multi-seed $|d|$ by $2.4\times$ at 50% contamination while reducing inter-seed coefficient of variation (CV) from 60% to 19%, establishing θ -anchoring as a necessary condition for reliable detection of high-ability models at high contamination. ARC-C reverses through a ceiling effect.

¹Code and data are available at https://anonymous.4open.science/r/irt_personfit_contam-6004/.




Figure 1: Overview of IRT person-fit detection. (a) A clean model’s responses track the IRT-predicted probability curve; (b) a contaminated model shows anomalous successes on hard items (red diamonds), producing a large $|\Delta \ell_z^*|$. (c) Effect-size landscape: baseline accuracy vs. Cohen’s $|d|$ across model families (diamonds: 5%, triangles: 15%, squares: 25%, circles: 50%; multi-seed means where available, single-seed otherwise; values from Table 1). (d) Dose-response trajectories; free- θ shows non-monotonic artifact, fixed- θ recovers consistent positive dose-response.

2. **Detection-diagnosis tradeoff.** Simpler statistics yield stronger detection: accuracy z -scores produce $|d|$ values several-fold to over an order of magnitude larger than person-fit statistics, and the nonparametric HT statistic (Van der Flier, 1982) generally exceeds ℓ_z^* . Yet the diagnostic power reverses: only ℓ_z^* supports seen/unseen decomposition, mechanistic θ -shift analysis, and sensitivity to response-pattern structure (57/60 vs. 4/60 flagged cells on real-world data).
3. **Non-monotonic dose-response under free- θ (exploratory).** The Qwen dose-response ($|d|$: 4.15 $\rightarrow$ 5.36 $\rightarrow$ 7.44 $\rightarrow$ 5.08 at 5%/15%/25%/50%; 5% is single-seed, remainder are multi-seed means) shows non-monotonicity; θ -anchoring eliminates the pat-

tern and reveals a broadly increasing dose-response ($|d|=12.32$ at 50%). The non-monotonicity is a θ -estimation artifact, not a detection boundary.

Our contributions are threefold: (1) We provide the first empirical characterization of IRT person-fit behavior for LLM contamination, quantifying theoretically predicted phenomena (θ -shift absorption, detection-diagnosis tradeoff) whose magnitudes and boundary conditions are not derivable from IRT theory alone. (2) We establish that $\Delta \ell_z^*$ reliably detects SFT-stage contamination across three model families (all single-seed bootstrap 95% CIs exclude zero, $n=10$). (3) We document the detection-diagnosis tradeoff, clarifying what binary response patterns can and cannot reveal. Findings 1–2 are confirmatory (pre-planned 15%/50%




Figure 2: Overview of IRT-FIT: (a) positioning in the detection frontier by information access; (b) pipeline from binary responses through ℓ_z^* person-fit analysis to detection and characterization; (c) three behavioral boundaries discovered; (d) recommended practitioner workflow.

× 3 families); Finding 3 is exploratory (5%/25% added post-hoc).

2 Related Work

Benchmark contamination detection. Contamination detection methods span a wide range of information requirements. Corpus-search methods match benchmark items against pretraining data (Dodge et al., 2021; Magar and Schwartz, 2022) but require full corpus access. Membership inference attacks use token-level log-probabilities (Shi et al., 2024; Deng et al., 2024); generation-based probes elicit benchmark content via prompting (Golchin and Surdeanu, 2024, 2025; Oren et al., 2024); performance-based methods compare accuracy across target and reference sets (Dekoninck et al., 2024; Zhang et al., 2024; Zawalski et al., 2026). Several works address contamination measurement and mitigation (Zhou et al., 2023; Sainz et al., 2023; Jacovi et al., 2023), while Bordt et al. (2025) show contamination can be “forgotten” under continued training, Han et al. (2026) extend contamination quantification to psychometric evaluations, and Kattamuri et al. (2025) use activation patterns for mechanistic detection. Fu et al. (2025) systematically evaluate detection method assumptions, finding degradation under realistic conditions.

A critical limitation unites token-level methods: Wang et al. (2025) show RL post-training renders membership inference attack (MIA)-based detectors no better than random, because RL reshapes token distributions while preserving binary correctness patterns. Tao et al. (2026) and Zhang et al. (2026) propose RL-specific solutions but require generation-level access. Our approach requires only binary response patterns and externally calibrated item parameters from PSN-IRT (Zhou et al., 2026).

IRT for LLM evaluation. IRT has seen growing adoption in ML evaluation: Martínez-Plumed et al. (2019) applied IRT to analyze classifiers at the instance level; Lalor et al. (2019) estimated IRT parameters from deep neural network (DNN) response patterns; Rodriguez et al. (2021) showed IRT-informed item weighting improves rankings; Polo et al. (2024) reduced benchmark sizes by 100× via IRT-based item selection; and Li et al. (2025) developed adaptive testing for efficient LLM evaluation. Most directly relevant, Zhou et al. (2026) calibrated 4PL item parameters across 11 benchmarks (41,871 items), providing the item bank our framework builds on. None of these works apply *person-fit* statistics for contamination detection.

Person-fit statistics in psychometrics. Person-fit analysis has a long history in educational measurement (see Embretson and Reise, 2000). Drasgow et al. (1985) introduced ℓ_z , a standardized log-likelihood residual; Snijders (2001) derived the corrected ℓ_z^* accounting for ability estimation uncertainty. Karabatsos (2003) compared 36 person-fit statistics and found ℓ_z among the most powerful for detecting cheating; we select ℓ_z^* for its theoretical grounding and HT (Van der Flier, 1982) as a nonparametric complement. Sinharay (2017) showed ℓ_z^* is particularly effective when compromised items are known, analogous to our seen/unseen decomposition. Our work transfers this psychometric toolkit to LLM evaluation, treating each model as an examinee whose response pattern may reveal memorization.

3 Method

3.1 IRT Person-Fit Framework

We adopt the four-parameter logistic (4PL) IRT model to characterize contamination-induced aberrance in LLM response patterns. The 4PL model defines the probability that a model with ability θ answers item i correctly:

$$P_i(\theta) = c_i + (d_i - c_i) \frac{1}{1 + e^{-a_i(\theta - b_i)}} \quad (1)$$

where a_i is discrimination, b_i difficulty, c_i the lower asymptote (guessing), and d_i the upper asymptote. The 4PL extension is essential for LLM evaluation: Zhou et al. (2026) find 46% of MMLU items have $d_i < 0.95$, meaning even high-ability models err on these items due to ambiguity or format effects; a 3PL model that forces $d_i=1$ misattributes such errors to low ability, inflating ℓ_z^* variance and reducing detection power (+55.7pp gain with 4PL; §4.3). Item parameters are pre-calibrated by PSN-IRT (Zhou et al., 2026) from a large-scale leaderboard dataset spanning 39 models; we use these external parameters rather than re-estimating from potentially contaminated data, since our marginal maximum likelihood EM (MML-EM) refit shows discrimination recovery degrades to $r_a=0.127$ on contaminated leaderboard data (Appendix I).

Given item parameters and a binary response vector $\mathbf{x} = (x_1, \dots, x_n)$, we compute the person-fit statistic ℓ_z^* (Snijders, 2001), a standardized log-likelihood residual quantifying deviation between observed and expected response patterns. The Snijders correction accounts for estimation uncertainty in $\hat{\theta}_{MLE}$, so that $\ell_z^* \sim \mathcal{N}(0, 1)$ under the




Figure 3: Dose-response trajectories on MMLU. Free- θ (solid) shows non-monotonic behavior for Qwen due to θ -shift absorption (shaded); fixed- θ (dashed) removes the inverted-U artifact and yields an overall positive dose-response ($|d|=12.32$ at 50%). Llama’s flat trajectory reflects minimal absorption. Hollow markers: single-seed ($n=1$); error bars: $\pm SD$ ($n=3$).

null (Drasgow et al., 1985; Snijders, 2001); Magis et al. (2012) provide a detailed treatment of how response model selection and ability estimation affect ℓ_z^* . Large negative ℓ_z^* indicates aberrant responding (unexpected failures on easy items or unexpected successes on hard items); large positive values indicate over-consistent patterns.

We consider two estimation modes: *free- θ* , where $\hat{\theta}_{MLE}$ is estimated jointly from the (potentially contaminated) response vector, and *fixed- θ* (θ -anchoring), where $\hat{\theta}$ is held at the clean-baseline mean. Under free- θ , contamination inflates $\hat{\theta}$, which partially “explains away” anomalous successes and attenuates ℓ_z^* ; fixed- θ eliminates this absorption but requires a clean reference (§4.2).

3.2 $\Delta\ell_z^*$ Matched-Difference Design

LLM response patterns systematically violate IRT assumptions of local independence and unidimensionality. In our empirical data, 57 of 60 model-benchmark cells yield $|\ell_z^*| > 1.96$ even absent contamination, rendering absolute thresholds uninformative. We address this through a matched-difference design.

For controlled experiments with clean baselines, we compute:

$$\Delta\ell_z^* = \ell_{z,\text{contam}}^* - \bar{\ell}_{z,\text{clean}}^* \quad (2)$$

where $\bar{\ell}_{z,\text{clean}}^*$ averages over $n=10$ clean seeds. The shared model misfit cancels, isolating

contamination-specific aberrance. We quantify effect sizes as $d = (\ell_{z,\text{contam}}^* - \bar{\ell}_{z,\text{clean}}^*) / s_{\text{clean}}$, where s_{clean} is the SD of the $n=10$ clean-seed ℓ_z^* values; with a single contaminated observation per condition, d is a standardized departure rather than a two-sample effect size (Hedges’ $g = J \cdot d$, $J \approx 0.914$ for $n=10$; [Hedges, 1981](#)). We report percentile bootstrap 95% CIs (10,000 resamples) reflecting clean-seed variability only; bias-corrected and accelerated (BCa) correction ([Appendix H](#)) generally confirms conservative coverage; multi-seed contamination-run variance appears in [Appendix L](#).

For deployment without clean baselines, cross-benchmark comparison provides the reference: the target benchmark’s ℓ_z^* is compared against benchmarks presumed clean from the same model. Significance thresholds derive from a nonparametric bootstrap null rather than the asymptotic $\mathcal{N}(0, 1)$ reference, accommodating residual model misspecification. Simulation analysis ([Appendix G.1](#)) confirms that the matched-difference design is robust to local dependence at realistic levels ($d_{\text{ratio}} \geq 0.93$ for $\sigma_\tau \leq 0.5$) and maintains $\geq 98.5\%$ power even under contamination-specific asymmetric dependence.

3.3 Experimental Setup

Models and contamination. We validate on three 7–8B model families: Qwen2.5-7B-Instruct, Llama-3.1-8B-Instruct, and Mistral-7B-Instruct-v0.3. Each is fine-tuned with Quantized Low-Rank Adaptation (QLoRA; rank = 16, $\alpha = 32$; full hyperparameters in [Appendix K](#)) under clean and contaminated conditions at varying levels: at $X\%$ contamination, $X\%$ of SFT training examples are benchmark items (questions paired with gold answers in the standard evaluation format), with the remainder drawn from a general instruction-following dataset. All three families are tested at 15% and 50% contamination. Following initial observation of non-monotonic dose-response for Qwen, we added 5% and 25% conditions for Qwen to trace the full dose-response curve. Both conditions were subsequently extended to Llama to test whether the pattern generalizes; Mistral was additionally tested at 5% but not at 25% due to computational constraints. We train $n=10$ independent clean seeds per model to establish baseline variability and construct bootstrap confidence intervals. Reported CIs reflect variability across n clean baseline seeds; for each model family, selected contamination levels were additionally replicated with

$n=3$ independent QLoRA seeds to assess training-run variance: Qwen at 15%/25%/50%, Llama at 5%/15%/25%/50%, and Mistral at 5%/15%/50% ([Appendix L](#)).

Benchmarks. We evaluate on MMLU ([Hendrycks et al., 2021](#)) (14,042 items) and ARC-Challenge ([Clark et al., 2018](#)) (295 PSN-IRT-calibrated items), spanning two orders of magnitude in benchmark size. All 4PL item parameters are from [Zhou et al. \(2026\)](#). ARC-Challenge experiments were conducted on Qwen only, as the primary evaluation of cross-benchmark generalizability; extending to additional families was constrained by computational budget.

Post-RL extension. To probe behavior boundaries beyond SFT, we apply Group Relative Policy Optimization (GRPO; [Shao et al., 2024](#)) to one contaminated Qwen model (15% MMLU) for 1,500 steps, evaluating at three checkpoints. This extension is $N=1$ (single model family, single contamination level).

Evaluation protocol. Ability estimation uses maximum likelihood (MLE) with Newton–Raphson optimization on the 4PL log-likelihood. For each model, we compute ℓ_z^* on the full benchmark response vector; seen/unseen decomposition recomputes ℓ_z^* on item subsets to localize the contamination signal. All evaluations use greedy decoding (temperature 0) with the standard multiple-choice prompt format from the lm-evaluation-harness ([Gao et al., 2023](#)).

Contamination model. QLoRA SFT on benchmark items proxies SFT-stage contamination, analogous to a test-taker with item preknowledge. At each contamination level, contaminated items are shuffled into the training mixture and presented in the same question–answer format used during evaluation, maximizing the memorization signal. We term contaminated items *seen* and the remainder *unseen*, enabling person-fit decomposition (§4.3). Findings are scoped to SFT-based contamination.

4 Results

4.1 Detection Signal Across Model Families

[Table 1](#) presents Cohen’s d effect sizes for ℓ_z^* across three model families and two benchmarks. All single-seed bootstrap CIs exclude zero, confirming that $\Delta\ell_z^*$ reliably separates contaminated from clean models at both contamination levels.

Table 1: Cohen’s d for $\Delta\ell_z^*$ across model families ($n=10$ clean seeds). Single-seed conditions: percentile bootstrap 95% CI (10,000 resamples). [†]Multi-seed conditions: mean $\pm$ SD ($n=3$ independent runs). All single-seed CIs exclude zero.

Model	Bench.	d (5%) [CI]	d (15%) [CI]	d (25%) [CI]	d (50%) [CI]
Qwen [†]	MMLU	-4.15 [-9.9, -2.7]	-5.36 (± 1.10)	-7.44 (± 1.86)	-5.08 (± 3.02)
Llama [†]	MMLU	-2.57 (± 0.99)	-2.80 (± 0.60)	-2.56 (± 0.52)	-2.30 (± 0.30)
Mistral [†]	MMLU	-3.83* (± 1.42)	-3.37 (± 0.68)	—	-4.09 (± 0.53)
Qwen [†]	ARC-C	—	+3.16 [2.0, 9.5]	—	+5.30 [3.6, 15.2]

[†]Multi-seed conditions: means ($n=3$ independent runs) shown as $\pm$ SD; remaining conditions are single-seed with bootstrap CIs.
* $n=2$ seeds. Per-seed values in Appendix L.




Figure 4: Cohen’s $|d|$ heatmap (values from Table 1; multi-seed means where available). Blue intensity reflects effect-size magnitude. Grey cells: not tested.

On MMLU, effect sizes range from $|d|=2.30$ (Llama at 50%) to $|d|=7.44$ (Qwen at 25%), with Qwen at 5% ($|d|=4.15$) confirming detection at the lowest contamination level tested. Llama ($|d|=2.57/2.80/2.56/2.30$) and Mistral ($|d|=3.83/3.37/4.09$) show comparable magnitudes. On ARC-C (295 items), Qwen yields $|d|=3.16$ at 15% and $|d|=5.30$ at 50%, demonstrating that detection remains feasible even on small benchmarks when contamination is sufficiently concentrated. The positive sign on ARC-C (vs. negative on MMLU) reflects a ceiling effect: with 89.7% baseline accuracy, contamination pushes responses toward over-consistency rather than the unexpected-success pattern seen on MMLU.

Clean-model ℓ_z^* baselines diverge substantially: Qwen +8.68 (over-consistent), Llama -0.14 (near-expected), Mistral +4.02 (Appendix D); seed sensitivity analysis confirms these reflect systematic misfit rather than seed artifacts (Appendix C), making absolute thresholds uninformative. The $\Delta\ell_z^*$ design normalizes these divergent baselines (Figure 4).

4.2 Model-Specific θ -Shift Dynamics

Contamination inflates $\hat{\theta}_{MLE}$, which “explains away” anomalous successes on hard items; we term this θ -shift absorption. θ -shift magnitude is comparable across families ($\Delta\theta \approx 0.5$), yet fixed- θ am-

plification diverges markedly (Table 2; Figure 1c).

θ -shift absorption is the dominant source of detection variability for high-ability models at high contamination: (i) fixed- θ estimation amplifies Qwen’s multi-seed $|d|$ by $2.4\times$ at 50% ($5.08 \rightarrow 12.32$; Table 2), while Llama shows no amplification ($0.4\times$); (ii) θ -anchoring reduces Qwen’s inter-seed CV from 60% to 19%, confirming that training-run variance is driven by θ -shift stochasticity rather than true signal variance; (iii) the cross-model ranking under free- θ is seed-dependent at 50%, making θ -anchoring a necessary condition for reliable detection of high-ability models at high contamination.

ARC-C reverses this pattern: despite 89.7% baseline accuracy, contamination produces *larger* effect sizes ($|d|=3.16/5.30$). This ceiling effect arises because 295 items with high baseline accuracy leave minimal room for θ inflation; contamination manifests as anomalous successes on the few remaining hard items, producing disproportionately large ℓ_z^* shifts. Detection sensitivity thus depends on the interaction between model ability and benchmark characteristics, suggesting that practitioners should interpret effect sizes relative to the specific model–benchmark pairing rather than applying universal thresholds.

Seen vs. unseen decomposition. To localize the contamination signal, we recompute ℓ_z^* separately on seen and unseen item subsets. At 15% contamination, Qwen’s seen-item ℓ_z^* drops to +0.67 while unseen items remain at +3.91 (seed 42; clean baseline +8.68), confirming that the person-fit signal originates predominantly from memorized items. This decomposition is unique to ℓ_z^* among the statistics we evaluate: neither accuracy z -scores nor HT support item-subset analysis.

Fixed- θ estimation. Table 2 compares free- and fixed- θ Cohen’s d . The amplification ratio tracks baseline $|\ell_z^*|$ monotonically: Qwen ($\ell_z^*=+8.68$) shows $2.4\times$ at 50%, Mistral (+4.02) $1.7\times$, while Llama (-0.14) shows $0.4\times$ despite comparable θ -shift ($\Delta\theta=0.52$). Fixed- θ is thus a targeted remedy for models with substantial baseline misfit, not a universal one.

4.3 Detection-Diagnosis Tradeoff

Table 3 compares three statistics: ℓ_z^* (4PL parameters), HT (difficulty only), and accuracy z -scores (no item parameters). For pure detection, simpler statistics are more powerful: accuracy z -scores

Table 2: Free- θ vs. fixed- θ Cohen’s $|d|$ on MMLU. Qwen[§]: multi-seed mean ($n=3$); Llama and Mistral: seed-42 single run. Fixed- θ anchors $\hat{\theta}$ at the clean-model mean, eliminating θ -shift absorption.

Model	$ d_{\text{free}} $	$ d_{\text{fixed}} $	Ratio	Fixed 95% CI / SD
<i>5% contamination</i>				
Qwen	4.15	3.52	0.8×	[2.9, 5.7]
<i>15% contamination</i>				
Qwen [§]	5.36	6.85	1.3×	(±1.80)
Llama	2.11	0.58	0.3×	[0, 2.37] [†]
Mistral [‡]	2.64	2.69	1.0×	[1.53, 9.06]
<i>50% contamination</i>				
Qwen [§]	5.08	12.32	2.4×	(±2.38)
Llama	2.33	0.84	0.4×	[0.17, 2.97]
Mistral [‡]	3.50	5.80	1.7×	[3.87, 18.45]

[†]CI includes zero.

[§]Multi-seed mean ($n=3$); CI column shows $\pm$ SD of fixed- θ $|d|$.

[‡]Fixed- θ : 6 seeds (response vectors from seeds 4–7 excluded due to incomplete evaluation runs); free- θ uses all 10. d_{free} CIs omitted; multi-seed means in Table 1.

yield $|d|$ values up to an order of magnitude larger than person-fit statistics, because cross-seed accuracy variance is near-zero ($\text{SD} \approx 0.0003$ for Qwen clean seeds). These extreme effect sizes reflect the controlled setting where training-run variance is minimal; in deployment without matched clean baselines, accuracy z -scores would lack this stability. HT generally exceeds ℓ_z^* , since HT is not attenuated by the Snijders correction.

Yet the diagnostic power reverses. ℓ_z^* uniquely supports (a) *seen/unseen decomposition*, though item-level detection is near-random (area under the ROC curve [AUROC] ≈ 0.53 ; Appendix F); (b) *θ -shift mechanism* (§4.2); and (c) *structural sensitivity*: ℓ_z^* flags 57/60 leaderboard cells as aberrant (Figure 5) vs. 4/60 for accuracy z -scores (St-Onge et al., 2011). In practice, accuracy z -scores screen candidates; ℓ_z^* then characterizes whether the signal reflects θ -shift absorption or direct memorization. Within the IRT family, 4PL gains +55.7pp power over 3PL on MMLU at 15% contamination (96.9% vs. 41.2%; Appendix B), because 46% of MMLU items have upper asymptote $d < 0.95$; the advantage persists across benchmarks (HellaSwag +34.3pp, GSM8K +20.7pp).

4.4 Ability-Dependent Non-Monotonic Dose-Response

As an exploratory observation, the free- θ dose-response is non-monotonic for Qwen. Figure 3 plots the dose-response trajectory, combining the single-seed 5% estimate (hollow marker) with multi-seed means at 15%/25%/50% ($|d|$: 4.15 $\rightarrow$ 5.36 $\rightarrow$ 7.44 $\rightarrow$ 5.08), showing signal peaking at 25% before declining; inter-seed variance at 50% is high ($\text{CV}=60\%$, range [1.90, 7.91]). This non-monotonicity is a θ -estimation artifact: fixed- θ

Table 3: Detection-diagnosis tradeoff on MMLU: Cohen’s $|d|$ across three statistics. Qwen[†]: multi-seed means ($n=3$); Llama and Mistral: seed-42 single run with bootstrap CIs.

Model	Statistic	$ d $ (15%)	$ d $ (50%)
Qwen [†]	ℓ_z^*	5.36 (± 1.10)	5.08 (± 3.02)
	HT	5.83 (± 1.30)	5.29 (± 3.41)
	Acc. z	45.0 (± 2.5)	70.0 (± 1.1)
Llama	ℓ_z^*	2.11 [1.4, 4.8]	2.33 [1.6, 5.1]
	HT	2.45 [1.5, 4.3]	3.07 [1.9, 5.3]
	Acc. z	22.5 [13.6, 47.0]	32.8 [19.8, 68.6]
Mistral	ℓ_z^*	2.64 [‡]	3.50 [2.1, 9.4]
	HT	4.36 [1.2, 7.0]	5.89 [1.9, 9.4]
	Acc. z	7.81 [4.4, 14.4]	16.6 [9.7, 34.6]

[†]Per-seed values in Appendix L. Bootstrap CIs: 10,000 resamples over $n=10$ clean seeds. [‡]Seed-42 single run; multi-seed $|d|=3.37$ (± 0.68) in Table 1.

eliminates the inverted-U ($|d|=12.32$ at 50%, CV reduced to 19%). Llama’s flat trajectory ($|d| \approx 2.3$ –2.8) is consistent with its near-zero baseline ℓ_z^* : the divergence is fully explained by baseline ability differences. Simulation corroborates: two-sided testing achieves $>93\%$ power at all levels on MMLU (Appendix A). These results demonstrate that non-monotonic dose-response is not a fundamental limitation of person-fit statistics but a predictable consequence of free- θ estimation in high-ability models, and that θ -anchoring provides a principled correction.

5 Discussion

Detection vs. diagnosis. The detection-diagnosis tradeoff reveals a design tension absent from prior contamination detection work. Simpler statistics (accuracy z -scores, HT) detect contamination more powerfully because they are not attenuated by the Snijders correction or inflated by θ -shift absorption. Yet ℓ_z^* provides unique mechanistic insight invisible to non-IRT statistics: θ -shift analysis reveals *why* detection varies across models (baseline ability determines absorption magnitude), *seen/unseen decomposition* localizes the signal source, and 4PL modeling captures ceiling effects that 3PL misses (+55.7pp power on MMLU). Our contribution lies not in proposing a new detection method but in quantifying these theoretically predicted phenomena where magnitudes are not derivable from IRT theory alone.

Practical implications. For practitioners, we recommend a two-stage workflow: (i) screen with accuracy z -scores, which require no item parameters




Figure 5: ℓ_z^* heatmap for 12 leaderboard models across five benchmarks. 57/60 cells exceed $|\ell_z^*| > 1.96$, indicating pervasive IRT assumption violations even absent contamination. Details in Appendix E.

and produce the largest effect sizes; (ii) characterize flagged models with ℓ_z^* to determine whether the signal reflects θ -shift absorption (high-ability model, large $\Delta\theta$) or direct memorization (anomalous successes on hard items without commensurate θ inflation). θ -anchoring is necessary to stabilize detection for high-ability models (free- θ CV up to 60% vs. 19% under θ -fixed; §4.2). Universal baseline misfit (57/60 cells flagged; Figure 5) means absolute ℓ_z^* thresholds cannot be used directly; the $\Delta\ell_z^*$ matched-difference design or cross-benchmark comparison is required.

Robustness and complementarity. PSN-IRT provides pre-calibrated 4PL parameters (Zhou et al., 2026) that are robust to estimation noise: $\pm 10\%$ parameter perturbation produces $< 4\%$ bias in fixed- θ Cohen’s d (Appendix G.2), because the matched-difference design cancels shared parameter error. Requiring only binary response patterns and an external item bank, ℓ_z^* occupies a unique position at the minimal-access end of the detection frontier: it complements logit-based methods (Shi et al., 2024; Dekoninck et al., 2024) that fail after RL post-training (Wang et al., 2025). Our GRPO pilot illustrates this complementarity: Min-K% AUC drops to chance (0.493) after 500 RL steps, while the ℓ_z^* response-pattern gap persists at ~ 3.3 across all checkpoints (Figure 6). This structural invariance (RL reshapes token-level distribu-



Figure 6: Post-RL trajectory (Qwen, 15% MMLU). GRPO erases the accuracy gap (11.4pp $\rightarrow$ 2.1pp) while the ℓ_z^* gap persists (~ 3.3). $N=1$; Appendix J.

tions but preserves binary correctness anomalies) suggests that person-fit statistics may remain informative in settings where all probability-based detectors fail, though this $N=1$ observation requires multi-family validation.

6 Conclusion

We presented the first empirical characterization of IRT person-fit statistics for LLM benchmark contamination detection, evaluating ℓ_z^* across three 7–8B model families (Qwen, Llama, Mistral), two benchmarks (MMLU, ARC-C), and four contamination levels (5%–50%). Our controlled experiments reveal three behavioral boundaries: (i) θ -shift absorption, where contamination inflates ability estimates and attenuates person-fit signals in a model-dependent manner, necessitating θ -anchoring for high-ability models; (ii) a detection-diagnosis tradeoff, where simpler statistics detect more powerfully but only ℓ_z^* supports mechanistic analysis; and (iii) non-monotonic dose-response under free- θ estimation, which is a θ -estimation artifact eliminated by anchoring. These findings establish that binary response patterns, combined with externally calibrated item parameters, carry sufficient information for contamination detection, though diagnostic θ -anchoring additionally requires a clean baseline; detection and diagnosis require different tools.

A preliminary post-RL pilot suggests person-fit anomalies persist after RL, complementing logit-based methods that fail in this setting. Future work should extend to larger models, pretraining-stage contamination, and multi-family post-RL validation.

Limitations

Our controlled experiments cover three 7–8B model families with QLoRA fine-tuning; generalization to larger scales, mixture-of-experts architectures, full-parameter fine-tuning, and pretraining-stage contamination remains untested. Most contamination conditions use a single training run, and multi-seed replication (available for Qwen, Llama, and Mistral) reveals high seed dependence at 50% contamination (CV=60%), partially resolved by θ -anchoring; the post-RL characterization is $N=1$. The framework depends on pre-calibrated 4PL item parameters from PSN-IRT, whose calibration sample may include contaminated models, and we do not provide head-to-head comparison with logit-based methods (MIA, ConStat) that operate at a different information level.

Ethical Considerations

This work demonstrates how SFT-stage benchmark contamination can be detected and characterized using IRT person-fit statistics. While the contamination methodology we employ (fine-tuning on benchmark items) is well-established in prior work, we acknowledge the dual-use nature of systematic contamination studies: the same controlled framework could inform adversarial contamination strategies. We believe the detection and diagnostic contributions outweigh this risk, as they support benchmark integrity and more reliable model evaluation. All experiments use publicly available benchmarks under permissive licenses and open-weight models; no private data or human subjects are involved.

Reproducibility Statement

Code and response data are available in an anonymized repository.² Bootstrap confidence intervals use the percentile method with 10,000 resamples. All three model families use clean baseline seeds 0–9 ($n=10$ per family). Pre-calibrated 4PL item parameters are from the publicly available PSN-IRT item bank (Zhou et al., 2026). Training used PyTorch 2.8.0, Transformers 5.8.1, PEFT 0.19.1, and BitsAndBytes 0.49.2. Both benchmarks are publicly available under permissive licenses (MMLU: MIT; ARC-Challenge: CC-BY-SA).

²https://anonymous.4open.science/r/irt_personfit_contam-6004/

References

- Sebastian Bordt, Suraj Srinivas, Valentyn Boreiko, and Ulrike von Luxburg. 2025. [How much can we forget about data contamination?](#) In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*.
- Y. Chang, X. Wang, J. Wang, Y. Wu, L. Yang, K. Zhu, H. Chen, X. Yi, C. Wang, Y. Wang, W. Ye, Y. Zhang, Y. Chang, P. S. Yu, Q. Yang, and X. Xie. 2024. [A survey on evaluation of large language models](#). *ACM Transactions on Intelligent Systems and Technology*.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. Think you have solved question answering? try ARC, the AI2 reasoning challenge. *arXiv preprint arXiv:1803.05457*.
- J. Dekoninck, M. Müller, and M. Vechev. 2024. [ConStat: Performance-based contamination detection in large language models](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- C. Deng, Y. Zhao, X. Tang, M. Gerstein, and A. Cohan. 2024. [Investigating data contamination in modern benchmarks for large language models](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
- J. Dodge, M. Sap, A. Marasović, W. Agnew, G. Ilharco, D. Groeneveld, M. Mitchell, and M. Gardner. 2021. [Documenting large webtext corpora: A case study on the colossal clean crawled corpus](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- F. Drasgow, M. V. Levine, and E. A. Williams. 1985. [Appropriateness measurement with polychotomous item response models and standardized indices](#). *British Journal of Mathematical and Statistical Psychology*, 38(1):67–86.
- Susan E. Embretson and Steven P. Reise. 2000. *Item Response Theory for Psychologists*. Lawrence Erlbaum Associates, Mahwah, NJ.
- Yujuan Fu, Ozlem Uzuner, Meliha Yetisgen, and Fei Xia. 2025. [Does data contamination detection work \(well\) for LLMs? A survey and evaluation on detection assumptions](#). In *Findings of the Association for Computational Linguistics: NAACL 2025*.
- Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golding, Jeffrey Hsu, Alber Le Noac’h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ociepa, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, and 5 others. 2023. [A framework for few-shot language model evaluation](#).
- S. Golchin and M. Surdeanu. 2024. [Time travel in LLMs: Tracing data contamination in large language](#)

657	models. In <i>Proceedings of the 12th International Conference on Learning Representations (ICLR)</i> .	712
658		713
659	Shahriar Golchin and Mihai Surdeanu. 2025. Data contamination quiz: A tool to detect and estimate contamination in large language models. <i>Transactions of the Association for Computational Linguistics (TACL)</i> , 13:809–830.	714
660		715
661		716
662		
663		
664	Jongwook Han, Woojung Song, Jonggeun Lee, and Yohan Jo. 2026. Quantifying data contamination in psychometric evaluations of LLMs. In <i>Findings of the Association for Computational Linguistics: EACL 2026</i> .	717
665		718
666		719
667		720
668		721
669	Larry V. Hedges. 1981. Distribution theory for Glass’s estimator of effect size and related estimators. <i>Journal of Educational Statistics</i> , 6(2):107–128.	722
670		723
671		724
672		725
673	D. Hendrycks, C. Burns, S. Basart, A. Zou, M. Mazeika, D. Song, and J. Steinhardt. 2021. Measuring massive multitask language understanding. In <i>Proceedings of the 9th International Conference on Learning Representations (ICLR)</i> .	726
674		727
675		
676		
677	A. Jacovi, A. Caciularu, O. Goldman, and Y. Goldberg. 2023. Stop uploading test data in plain text: Practical strategies for mitigating data contamination by evaluation benchmarks. In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i> .	728
678		729
679		730
680		731
681		732
682		733
683	G. Karabatsos. 2003. Comparing the aberrant response detection performance of thirty-six person-fit statistics. <i>Applied Measurement in Education</i> , 16(4):277–298.	734
684		735
685		
686		
687	Ashish Kattamuri, Harshwardhan Fartale, Arpita Vats, Rahul Raja, and Ishita Prasad. 2025. RADAR: Mechanistic pathways for detecting data contamination in LLM evaluation. In <i>NeurIPS 2025 Workshop on Evaluating the Evolving LLM Lifecycle</i> .	736
688		737
689		738
690		739
691		740
692		741
693	John P. Lalor, Hao Wu, and Hong Yu. 2019. Learning latent parameters without human response patterns: Item response theory with artificial crowds. In <i>Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-IJCNLP)</i> , pages 4249–4259. Association for Computational Linguistics.	742
694		743
695		744
696		745
697		746
698		747
699		
700	P. Li, X. Tang, S. Chen, Y. Cheng, R. Metoyer, T. Hua, and N. V. Chawla. 2025. Adaptive testing for LLM evaluation: A psychometric alternative to static benchmarks. <i>arXiv preprint arXiv:2511.04689</i> .	748
701		749
702		750
703		751
704	I. Magar and R. Schwartz. 2022. Data contamination: From memorization to exploitation. In <i>Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)</i> .	752
705		753
706		754
707		755
708	David Magis, Gilles Raïche, and Sébastien Béland. 2012. A didactic presentation of Snijders’s ℓ_z^* index of person fit with emphasis on response model selection and ability estimation. <i>Journal of Educational and Behavioral Statistics</i> , 37(1):57–81.	756
709		757
710		758
711		759
	F. Martínez-Plumed, R. B. C. Prudêncio, A. Martínez-Usó, and J. Hernández-Orallo. 2019. Item response theory in AI: Analysing machine learning classifiers at the instance level. <i>Artificial Intelligence</i> , 271:18–42.	760
		761
		762
	Y. Oren, N. Meister, N. S. Chatterji, F. Ladhak, and T. Hashimoto. 2024. Proving test set contamination in black box language models. In <i>Proceedings of the 12th International Conference on Learning Representations (ICLR)</i> .	763
		764
		765
		766
		767
	F. M. Polo, L. Weber, L. Choshen, Y. Sun, G. Xu, and M. Yurochkin. 2024. tinybenchmarks: Evaluating LLMs with fewer examples. In <i>Proceedings of the 41st International Conference on Machine Learning (ICML)</i> , volume 235 of <i>Proceedings of Machine Learning Research</i> , pages 34303–34326. PMLR.	
	Pedro Rodriguez, Joe Barrow, Alexander Hoyle, John P. Lalor, Robin Jia, and Jordan Boyd-Graber. 2021. Evaluation examples are not equally informative: How should that change NLP leaderboards? In <i>Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)</i> , pages 4486–4503. Association for Computational Linguistics.	
	O. Sainz, J. Campos, I. García-Ferrero, J. Etxaniz, O. Lopez de Lacalle, and E. Agirre. 2023. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In <i>Findings of the Association for Computational Linguistics: EMNLP 2023</i> .	
	Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y.K. Li, Y. Wu, and Daya Guo. 2024. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. <i>arXiv preprint arXiv:2402.03300</i> .	
	W. Shi, A. Ajith, M. Xia, Y. Huang, D. Liu, T. Blevins, D. Chen, and L. Zettlemoyer. 2024. Detecting pre-training data from large language models. In <i>Proceedings of the 12th International Conference on Learning Representations (ICLR)</i> .	
	Sandip Sinharay. 2016. Assessment of person fit using resampling-based approaches. <i>Journal of Educational Measurement</i> , 53(1):63–85.	
	Sandip Sinharay. 2017. Which statistic should be used to detect item preknowledge when the set of compromised items is known? <i>Applied Psychological Measurement</i> , 41(6):403–421.	
	T. A. B. Snijders. 2001. Asymptotic null distribution of person fit statistics with estimated person parameter. <i>Psychometrika</i> , 66(3):331–342.	
	Christina St-Onge, Pierre Valois, Belkacem Abdous, and Stéphane Germain. 2011. Accuracy of person-fit statistics: A Monte Carlo study of the influence of aberrance rates. <i>Applied Psychological Measurement</i> , 35(6):419–432.	

Y. Tao, T. Wang, Y. Dong, H. Liu, K. Zhang, X. Hu, and G. Li. 2026. [Detecting data contamination from reinforcement learning post-training for large language models](#). In *Proceedings of the 14th International Conference on Learning Representations (ICLR)*.

H. Van der Flier. 1982. [Deviant response patterns and comparability of test scores](#). *Journal of Cross-Cultural Psychology*, 13(3):267–298.

H. Wang, H. Li, B. Ko, and H. Zhang. 2025. [On the fragility of benchmark contamination detection in reasoning models](#). *arXiv preprint arXiv:2510.02386*.

M. Zawalski, M. Boudir, K. Bałazy, B. Nushi, and P. Ribalta. 2026. [Detecting data contamination in LLMs via in-context learning](#). In *Proceedings of the 14th International Conference on Learning Representations (ICLR)*.

H. Zhang, Y. Yang, J. Yan, G. Bao, Y. Zhang, and Y. Zhang. 2026. [Detecting RLVR training data via structural convergence of reasoning](#). *arXiv preprint arXiv:2602.11792*.

Huixuan Zhang, Yun Lin, and Xiaojun Wan. 2024. [PaCoST: Paired confidence significance testing for benchmark contamination detection in large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*.

H. Zhou, H. Huang, Z. Zhao, L. Han, H. Wang, K. Chen, M. Yang, W. Bao, J. Dong, B. Xu, C. Zhu, H. Cao, and T. Zhao. 2026. [Lost in benchmarks? rethinking large language model benchmarking with item response theory](#). In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, volume 40, pages 35085–35093.

Kun Zhou, Yutao Zhu, Zhipeng Chen, Wentong Chen, Wayne Xin Zhao, Xu Chen, Yankai Lin, Ji-Rong Wen, and Jiawei Han. 2023. [Don’t make your LLM an evaluation benchmark cheater](#). *arXiv preprint arXiv:2311.01964*.

A Simulation Validation

Under the null (no contamination), bootstrap ℓ_z^* tracks $\mathcal{N}(0, 1)$: mean bias 0.041 and standard deviation 1.003 on MMLU. Kolmogorov–Smirnov rejection rates vary with benchmark size: 8.3% for MMLU and HellaSwag ($n > 10,000$), 50.0% for GSM8K ($n = 1,319$), 58.3% for ARC-C ($n = 295$), and 75.0% for HumanEval ($n = 164$). ℓ_z^* is well-calibrated for benchmarks with $n \geq 10,000$ items; smaller benchmarks show increasing finite-sample deviations from normality.

Two-sided testing eliminates the inverted-U artifact: MMLU achieves 97.8% power at 15% contamination and sustains $>93\%$ at all levels. Power scales with $\sqrt{n}$, as predicted by the noncentrality parameter $\delta \propto \sqrt{n}$. The empirical power function,

derived via Gauss–Hermite quadrature ($K = 30$) over the ability distribution, achieves $R^2 = 0.957$ against simulated values, with approximate minimum benchmark size of $n \geq 7,958$ items for 80% power at 15% contamination under MMLU item parameters.

B 4PL vs. 3PL Ablation Details

Table 4 reports detection power (proportion of 2,000 simulation replicates rejecting H_0 at $\alpha=0.05$, two-sided) for 4PL and 3PL parameterizations across five benchmarks and seven contamination levels. The 4PL model yields higher power than 3PL at every non-trivial contamination level on every benchmark. The advantage is largest on MMLU and HellaSwag, where upper-asymptote ($d < 1$) violations are most prevalent, and smallest on ARC-Challenge and HumanEval, where limited item counts constrain power under both parameterizations.

Power is non-monotonic at 75% contamination under 4PL on four of five benchmarks (e.g., MMLU: .952 $\rightarrow$.938; GSM8K: .846 $\rightarrow$.802). This reflects the θ -shift boundary condition described in §4.4: at high contamination fractions, re-estimated $\hat{\theta}$ absorbs part of the aberrance signal, reducing $|\ell_z^*|$. At 100% contamination, all items are seen and power reaches 1.000 trivially. The 4PL advantage over 3PL depends on accurate estimation of the upper asymptote d ; with estimation error exceeding $\sim 10\%$, the effective model behavior approaches 3PL ($d=1$) and the power gain diminishes. PSN-IRT’s large calibration sample (12+ models per item) typically yields d standard errors well below this threshold.

C Seed Sensitivity Analysis

Clean-baseline variability. Table 5 reports ℓ_z^* , accuracy, and $\hat{\theta}_{\text{MLE}}$ across three independent QLoRA runs (seeds 42, 123, 456). For Qwen2.5-7B, ℓ_z^* ranges from 7.05 (seed 123) to 9.53 (seed 456); the lower value of seed 123 likely reflects seed-specific training dynamics (random initialization, data ordering), but all three seeds produce $\ell_z^* \gg 1.96$, confirming that the elevated baseline reflects systematic model–benchmark misfit rather than seed-specific artifacts. Llama-3.1-8B shows ℓ_z^* near zero (-0.14 ± 0.97), consistent with lower baseline misfit. Accuracy variability is negligible for both families (CV < 0.004).

Table 4: Detection power (proportion rejecting H_0 at $\alpha=0.05$) under 4PL and 3PL parameterizations across benchmarks and contamination levels ($n=2,000$ replicates per cell). Bold indicates the 4PL advantage exceeds 5 pp.

Benchmark	IRT	5%	10%	15%	25%	50%	75%	100%
MMLU	4PL	.476	.848	.969	.983	.952	.938	1.00
	3PL	.227	.340	.412	.521	.676	.770	1.00
HellaSwag	4PL	.251	.454	.656	.915	.939	.919	1.00
	3PL	.167	.241	.313	.368	.508	.642	1.00
GSM8K	4PL	.156	.309	.492	.737	.846	.802	1.00
	3PL	.102	.172	.285	.396	.567	.577	1.00
ARC-C	4PL	.057	.075	.087	.130	.285	.493	1.00
	3PL	.054	.062	.071	.076	.110	.102	1.00
HumanEval	4PL	.076	.121	.197	.302	.554	.521	1.00
	3PL	.063	.086	.120	.182	.281	.315	1.00

Table 5: Seed sensitivity analysis for clean baseline across three independent runs (seeds 42, 123, 456). CV = coefficient of variation (std/|mean|). These $n=3$ pilot results are superseded by the $n=10$ analysis in §4.1; retained for methodological transparency.

Model	Metric	Mean	Std	Range	CV
Qwen2.5-7B	ℓ_z^*	8.684	1.412	2.475	0.163
	Accuracy	0.679	0.000	0.001	0.001
	$\hat{\theta}_{MLE}$	-0.097	0.012	0.023	0.119
Llama-3.1-8B	ℓ_z^*	-0.141	0.965	1.929	6.860
	Accuracy	0.594	0.002	0.004	0.004
	$\hat{\theta}_{MLE}$	-0.645	0.019	0.036	0.029

Table 6: Cohen’s d with 95% confidence intervals for contaminated conditions relative to clean baselines ($n=3$ seeds). Analytic CIs use $t_{0.025, df=2}$; bootstrap CIs use 10,000 resamples. These $n=3$ pilot results are superseded by the $n=10$ analysis in §4.1; retained for methodological transparency.

Model	Cond.	d	Analytic CI	Bootstrap CI
Qwen	15%	-3.42	[-9.92, 3.08]	[-3.45, -2.82]
	50%	-0.99	[-4.02, 2.04]	[-0.99, -0.41]
Llama	15%	-2.95	[-8.70, 2.80]	[-5.89, -2.29]
	50%	-3.31	[-9.63, 3.01]	[-6.47, -2.59]

Cohen’s d confidence intervals. Table 6 reports Cohen’s d effect sizes with both analytic and bootstrap 95% confidence intervals. Analytic CIs use $d \pm t_{0.025, df=2} \cdot SE_d$; with only $n=3$ clean seeds ($df=2$), the critical value $t_{0.025, 2} = 4.303$ produces intervals wide enough to cross zero in all four conditions. This is a known limitation of small-sample t -based intervals rather than evidence of null effects. Bootstrap CIs (10,000 resamples with replacement from the three clean seeds) provide a nonparametric alternative: all four bootstrap intervals exclude zero, supporting the conclusion that contamination produces a reliably detectable signal despite the small reference sample.

Cross-family consistency. Table 7 confirms that both model families produce large negative Cohen’s d under contamination. The $|\Delta d| = 2.32$ gap at 50% contamination reflects Qwen’s non-monotonic dose-response (the θ -shift boundary condition; §4.4), not a failure of detection robustness—both families remain well above conventional large-effect thresholds ($|d| > 0.8$) at both contamination levels.

Table 7: Cross-family consistency of contamination detection ($n=3$ pilot). Both families show large negative d under contamination. Superseded by $n=10$ results in §4.1; retained for methodological transparency.

Condition	Qwen d	Llama d	$ \Delta d $	Consistent?
15% contam.	-3.42	-2.95	0.47	Yes
50% contam.	-0.99	-3.31	2.32	Yes

D Baseline ℓ_z^* Interpretation

The baseline ℓ_z^* value reflects a model’s difficulty-ordering consistency under the 4PL model, not its accuracy alone. In the PSN-IRT leaderboard data (12 models $\times$ 5 benchmarks, distinct from our controlled experiments), Mistral (acc=0.514, lowest among the 12 models) exhibits $\ell_z^* = +5.42$ (second highest), because its response pattern follows the IRT-predicted difficulty ordering more closely than higher-accuracy models whose responses violate the assumed item structure. This dissociation explains why Mistral produces the strongest contamination signal despite the lowest accuracy: its lower baseline ability ($\hat{\theta} = -1.16$) provides greater θ -shift headroom.

Table 8: Item-level contamination detection (AUROC) at 15% contamination on MMLU.

Method	AUROC
Random baseline	0.50
ℓ_z^* itemwise	0.53
LR cross-model	0.59
RF cross-model	0.63

E Leaderboard Aberrance Analysis

We computed ℓ_z^* and accuracy z -scores for all 12 models $\times$ 5 benchmarks = 60 cells in the PSN-IRT leaderboard dataset. Using the standard $|\ell_z^*| > 1.96$ threshold ($\alpha=0.05$, two-tailed), 57 of 60 cells are flagged as aberrant, indicating pervasive IRT assumption violations in production LLM responses. In contrast, accuracy z -scores ($|z| > 1.96$) flag only 4 of 60 cells. This stark asymmetry—95% aberrance by ℓ_z^* versus 7% by accuracy—motivates the $\Delta\ell_z^*$ matched-difference design (§3.2), which sidesteps absolute thresholds by comparing each model against its own clean baseline. Figure 5 visualizes this pattern.

F Item-Level Detection Results

Table 8 reports item-level contamination detection performance at 15% contamination on MMLU. All methods perform near random (AUROC ≤ 0.63), confirming that contamination is diffuse across the response pattern and that model-level aggregation via ℓ_z^* is essential.

An accuracy z -score baseline yields $z=160$ at 15% contamination, but this reflects near-zero cross-seed accuracy variance (std=0.0003), not genuine discriminability. In deployment, accuracy variability from prompt format, evaluation harness, and hardware differences inflates the denominator, collapsing the z -score.

G Robustness Analysis

G.1 $\Delta\ell_z^*$ Under Local Dependence

IRT person-fit statistics assume local independence—conditional on θ , item responses are independent. LLM benchmarks violate this assumption through topical clustering (e.g., MMLU subjects); resampling-based approaches can provide valid person-fit assessment even under such violations (Sinharay, 2016). We evaluate whether the $\Delta\ell_z^*$ matched-difference design is robust to local dependence (LD) via two simulation

Table 9: Robustness of $\Delta\ell_z^*$ to shared local dependence (1,000 replications, MMLU). d_{ratio} : effect size relative to the $\sigma_\tau=0$ baseline.

σ_τ	Abs. ℓ_z^* Type I	d_{ratio} (15%)	d_{ratio} (50%)
0.0	0.043	1.000	1.000
0.3	0.110	1.002	1.038
0.5	0.454	0.934	1.028
1.0	0.939	0.705	1.061

Table 10: $\Delta\ell_z^*$ detection under asymmetric local dependence ($\sigma_\tau=0.5$, contamination-specific memory factor σ_{mem} ; 1,000 replications, MMLU).

σ_{mem}	Power (15%)	d (15%)	Power (50%)	d (50%)
0.5	0.987	−2.107	0.991	−2.337
1.0	0.993	−2.104	0.985	−1.587

scenarios on 14,042 MMLU items with 1,000 replications each.

Shared LD. Clean and contaminated conditions share identical testlet structure (~ 200 testlets of ~ 70 items, random effect $\gamma_t \sim \mathcal{N}(0, \sigma_\tau)$). Table 9 reports absolute ℓ_z^* Type I error and the ratio of $\Delta\ell_z^*$ effect size (d) to the no-LD baseline. While absolute ℓ_z^* Type I error inflates severely under LD (0.043 $\rightarrow$ 0.939 at $\sigma=1.0$), the $\Delta\ell_z^*$ design maintains detection power: $d_{\text{ratio}} \geq 0.93$ for $\sigma_\tau \leq 0.5$ (realistic range), confirming that shared misfit is cancelled by the matched-difference design.

Asymmetric LD. In the second scenario, contaminated models’ seen items share an additional memory factor ($\gamma_{\text{mem}} \sim \mathcal{N}(0, \sigma_{\text{mem}})$), simulating contamination-specific LD not present in clean conditions ($\sigma_\tau=0.5$ fixed). Table 10 shows that $\Delta\ell_z^*$ maintains $\geq 98.5\%$ power across all conditions, even when contamination introduces asymmetric dependence among seen items.

G.2 Sensitivity to Item Parameter Estimation Error

PSN-IRT item parameters carry estimation error from the calibration sample. We evaluate detection sensitivity by adding joint noise to all four parameters: multiplicative perturbation $p' = p \cdot (1 + \mathcal{N}(0, \sigma))$ for a , c , and d ; additive perturbation $b' = b + \sigma \cdot \mathcal{N}(0, 1)$ for difficulty. Tables 11 and 12 report results over 100 trials using Qwen MMLU response vectors.

Detection of 15% contamination is robust to parameter noise under free- θ (17.4% bias at $\sigma=10\%$, $|d|$ still >5 ; Table 11). The 50% contamination

Table 11: Sensitivity of detection to item parameter noise (100 trials, Qwen MMLU, free- θ). % bias relative to noise-free d .

Noise σ	d (15%)	% bias	d (50%)	% bias
0%	-6.575	—	-1.897	—
5%	-6.010	8.6	-0.925	51.2
10%	-5.433	17.4	-0.014	99.3

Table 12: Sensitivity of detection to item parameter noise (100 trials, Qwen MMLU, θ -fixed). Worst-case $|d|$ over 100 trials in parentheses.

Noise σ	d_{fixed} (15%)	% bias	d_{fixed} (50%)	% bias
0%	4.77	—	14.52	—
5%	5.07 ± 0.24	6.3	14.76 ± 0.44	1.6
10%	4.95 ± 0.43 (3.73)	3.7	14.45 ± 0.85 (12.20)	0.5

condition under free- θ is fragile—consistent with its already-weak signal due to θ -shift absorption (§4.4).

Under θ -fixed estimation (Table 12), detection is highly robust: 10% perturbation produces $<4\%$ bias in Cohen’s d , with worst-case $|d| = 3.73$ (15% contamination) and 12.20 (50%)—both far exceeding conventional thresholds. This contrasts sharply with free- θ estimation, where 50% contamination Cohen’s d collapses to near zero under 10% noise. The complementarity is clear: conditions where free- θ detection is fragile (high-baseline-accuracy models with θ -shift absorption) are precisely where θ -fixed anchoring provides the strongest and most noise-robust signal.

H BCa Bootstrap Sensitivity Check

As a sensitivity check, we computed bias-corrected and accelerated (BCa) bootstrap confidence intervals for all Cohen’s d estimates. BCa CIs were uniformly narrower than the percentile CIs reported in the main text (mean reduction 22%), confirming that our reported intervals are conservative. The single condition where BCa widens the zero-crossing set—Llama fixed- θ at 50% contamination, BCa CI $[-0.13, 2.21]$ vs. percentile CI $[0.17, 2.97]$ —is already characterized as ineffective in §4.2.

I Classical IRT Refit from Leaderboard Data

To assess whether item parameters can be recovered without PSN-IRT’s neural calibration, we applied classical MML-EM estimation (Gauss–Hermite quadrature, $Q=31$ nodes) to binary response data from $N=178$ models on the Open

Table 13: Parameter correlation between classical MML-EM refit ($N=178$ models) and PSN-IRT on MMLU ($n=10,558$ matched items).

Model	Parameter	Pearson r	Spearman ρ
3PL	a (discrim.)	0.127	0.135
3PL	b (difficulty)	0.663	0.684
2PL	a (discrim.)	-0.029	0.050
2PL	b (difficulty)	0.608	0.656

LLM Leaderboard v1. With 14,042 MMLU items and an examinee-to-item ratio of $N:J = 0.013$ —far below the $N>500$ typically required for stable 3PL estimation—discrimination parameters are severely under-determined.

Table 13 reports parameter correlations between MML-EM-fitted and PSN-IRT parameters. Difficulty shows moderate agreement ($r_b=0.663$), but discrimination is essentially unrecovered ($r_a=0.127$). Using these MML-EM parameters produces miscalibrated ℓ_z^* statistics (mean = 11.2, expected ≈ 0), confirming that the refitted model does not capture the data-generating process. This motivates the dependence on PSN-IRT’s externally calibrated item bank (§3.1).

J Post-RL Preliminary Observation

Wang et al. (2025) show that all MIA-based detectors fail after RL post-training. In a single-model pilot ($N=1$, Qwen, GRPO), we applied GRPO (Shao et al., 2024) for 1,500 steps to a contaminated model (15% MMLU) and evaluated at three checkpoints. GRPO erases the accuracy gap between seen and unseen items (11.4pp $\rightarrow$ 2.1pp), but the ℓ_z^* response-pattern gap persists (~ 3.3 ; Figure 6), suggesting structural traces of contamination may survive RL alignment. Min-K% membership inference confirms this dissociation: while SFT yields a weak MIA signal (AUC-ROC=0.611), GRPO training drives AUC to 0.493 (step 500), 0.493 (step 1000), and 0.493 (step 1500)—indistinguishable from random chance (0.50) as RL erases memorization-based features while the response-pattern anomaly that ℓ_z^* detects persists. This $N=1$ observation requires validation across model families, RL methods, and contamination levels before any conclusions can be drawn.

GRPO hyperparameters. Table 14 reports the GRPO training hyperparameters for the post-RL extension (§3.3).

Table 14: GRPO training hyperparameters (Qwen2.5-7B-Instruct, 15% MMLU contamination).

Hyperparameter	Value
Learning rate	5×10^{-6}
Per-device batch size	2
Gradient accumulation steps	4
Num. generations per prompt	16
Max completion length	128
Sampling temperature	1.2
Training steps	1,500 (checkpoints at 500, 1000, 1500)
Reward function	Correctness-based (+1.0 correct, -1.0 incorrect)
LoRA rank / α / dropout	16 / 32 / 0.05
Precision	BF16
Seed	42

Table 15: QLoRA fine-tuning hyperparameters (shared across all model families and contamination conditions).

Hyperparameter	Value
Quantization	4-bit NF4
LoRA rank	16
LoRA α	32
LoRA dropout	0.05
Target modules	q_proj, k_proj, v_proj, o_proj
Learning rate	2×10^{-4}
Optimizer	AdamW
Warmup ratio	0.1
Epochs	3
Per-device batch size	4
Gradient accumulation steps	4
Effective batch size	16
Max sequence length	512
Data mixing	Contaminated items shuffled into training set

K QLoRA Training Hyperparameters

Table 15 reports the QLoRA fine-tuning hyperparameters used across all experiments. All training was conducted on NVIDIA RTX PRO 6000 GPUs (48 GB VRAM). Each QLoRA SFT run takes approximately 15–25 minutes depending on contamination level and dataset size; the full experimental campaign (including all seeds, contamination levels, and model families) consumed approximately 120 GPU-hours.

L Contamination Run Reproducibility

Tables 16–19 report per-seed Cohen’s d for all model families across contamination levels, computed against the shared $n=10$ clean baseline. These tables underpin the multi-seed means and SDs reported in Tables 1 and 3.

M Use of AI Assistance

We used Claude for coding, code drafting, debugging, and language polishing.

Table 16: Reproducibility of contamination effect sizes across independent QLoRA runs (Qwen2.5-7B-Instruct, MMLU). Each row shows Cohen’s $|d|$ from an independent training seed against the shared $n=10$ clean baseline ($\ell_z^* = 8.68$, $s = 0.74$). All runs use identical hyperparameters; only the random seed differs.

Seed	Free- θ $ d $			Fixed- θ $ d $		
	15%	25%	50%	15%	25%	50%
42 (original)	6.57	8.22	1.90	4.77	4.96	14.52
123	5.09	8.78	5.42	7.79	5.37	12.64
456	4.41	5.31	7.91	7.99	7.34	9.80
Mean	5.36	7.44	5.08	6.85	5.89	12.32
SD	1.10	1.86	3.02	1.80	1.27	2.38
CV (%)	20.3	25.2	59.6	26.3	21.7	19.3

Note. Fixed- θ estimation anchors $\hat{\theta}$ at the clean-baseline mean ($\theta = -0.102$), removing θ -shift absorption. The 50% free- θ column exhibits high CV (59.6%), reflecting seed-dependent θ -shift magnitude; fixed- θ reduces this to 19.3%, confirming that cross-seed variability at high contamination is driven by ability-estimate instability rather than true signal variance.

Table 17: Multi-seed Cohen’s $|d|$ for HT and accuracy z -score statistics (Qwen2.5-7B-Instruct, MMLU), computed using the chi-squared ratio HT statistic and bootstrap Cohen’s d (10,000 resamples). These supplement the ℓ_z^* values in Table 3.

Seed	HT $ d $			Acc. z $ d $		
	15%	25%	50%	15%	25%	50%
42 (original)	7.26	8.90	1.66	42.3	49.8	68.7
123	5.52	9.51	5.80	47.3	53.4	70.8
456	4.72	5.10	8.41	45.4	51.3	70.4
Mean	5.83	7.84	5.29	45.0	51.5	70.0
SD	1.30	2.39	3.41	2.5	1.8	1.1

Note. HT exhibits high cross-seed variance at 50% (SD=3.41), mirroring the ℓ_z^* pattern: the non-monotonicity observed in seed 42 ($|d|=1.66$) does not generalize across seeds. Accuracy z -scores are highly stable (CV < 6%) and monotonically increasing with contamination level.

Table 18: Multi-seed Cohen’s d for Llama-3.1-8B across all contamination levels (MMLU). Each cell shows Cohen’s d from an independent training seed against the shared $n=10$ clean baseline ($\bar{\ell}_z^* = +0.26$, $s = 1.54$). The 5% condition exhibits higher cross-seed variance (SD=0.99), reflecting the weaker signal from ~ 700 contaminated items ($5\% \times 14,042$).

Contam. level	Per-seed d			Mean	SD
	s_1	s_2	s_3		
5%	−3.22	−3.06	−1.44	−2.57	0.99
15%	−2.11	−3.21	−3.09	−2.80	0.60
25%	−3.16	−2.28	−2.24	−2.56	0.52
50%	−2.33	−2.59	−1.99	−2.30	0.30

Note. Unlike Qwen, Llama exhibits no non-monotonic dose-response: mean $|d|$ is approximately flat across contamination levels (2.30–2.80), consistent with its near-zero baseline ℓ_z^* (+0.26) and negligible θ -shift absorption (§4.2). Cross-seed variance decreases with contamination level (CV: 39% at 5% $\rightarrow$ 13% at 50%), suggesting that higher contamination produces a more stable signal despite the flat mean trajectory.

Table 19: Multi-seed Cohen’s d for Mistral-7B-v0.3 across contamination levels (MMLU). Each cell shows Cohen’s d from an independent training seed against the shared $n=10$ clean baseline ($\bar{\ell}_z^* = +4.02$, $s = 2.05$). The 5% condition has only $n=2$ seeds; its high SD reflects this limited sample.

Contam. level	Per-seed d			Mean	SD
	s_1	s_2	s_3		
5%	−2.82	−4.83	—	−3.83	1.42
15%	−2.64	−3.98	−3.48	−3.37	0.68
50%	−3.50	−4.25	−4.51	−4.09	0.53